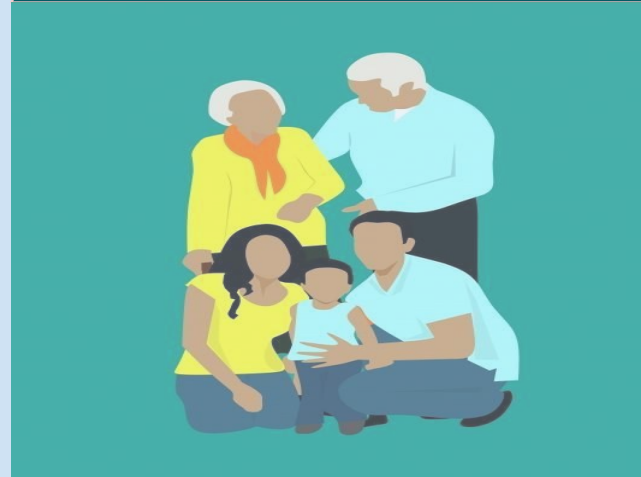


Home, Home Alone



Common Real Estate Problems





Leveraging ML to Better Inform Real Estate Agents

A photograph of two hands shaking in a firm grip, symbolizing a deal or agreement. The hands are wearing dark grey suit sleeves. In the background, a two-story house with light-colored siding and a dark roof is visible under a blue sky with some clouds. The entire image is framed by a teal background with a white diagonal line running from the top-left to the bottom-right.

Leveraging ML to Better Inform Potential Homeowners

Potential Impact

New Home Owners



Real Estate Agents



Potential Impact

New Home Owners



- Home-Price Fit
- Increased Utility
- Faster Closing Process

Potential Impact

- Increased Satisfaction
- Targeted Customer Outreach
- Faster Closing Process
- Efficient Use of Inventory



EFFICIENCY

**More Efficient Real
Estate Market**



Raw Dataset

- 5,484,743 Rows, 58 Total Columns
- Kaggle data set, updated frequently
- 02/01/2012 - 05/31/2024
- Redfin
- 5,215,6152 null values
- <https://www.senate.gov/senator/s/Senators1789toPresent.htm>
- Data Dictionary

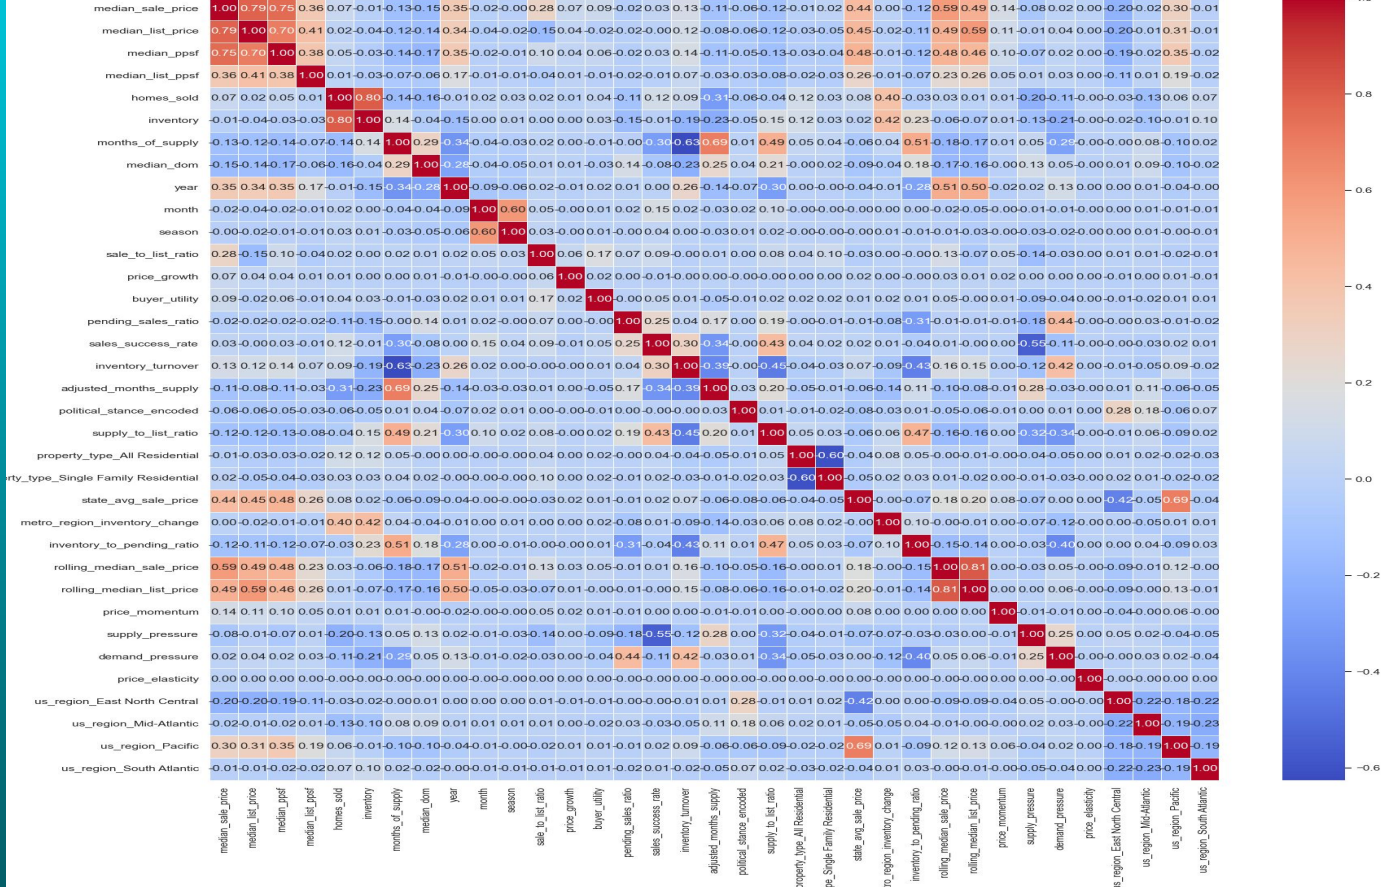
Processed Dataset

```
1 median_sale_price float64
2 median_list_price float64
3 median_ppsf float64
4 median_list_ppsf float64
5 homes_sold float64
6 inventory float64
7 months_of_supply float64
8 median_dom float64
9 year float64
10 month float64
11 season float64
12 sale_to_list_ratio float64
13 price_growth float64
14 buyer_utility float64
15 pending_sales_ratio float64
16 sales_success_rate float64
17 inventory_turnover float64
18 adjusted_months_supply float64
19 political_stance_encoded float64
20 supply_to_list_ratio float64
21 property_type_All Residential float64
22 property_type_Single Family Residential float64
23 state_avg_sale_price float64
24 metro_region_inventory_change float64
25 inventory_to_pending_ratio float64
26 rolling_median_sale_price float64
27 rolling_median_list_price float64
28 price_momentum float64
29 supply_pressure float64
30 demand_pressure float64
31 price_elasticity float64
32 us_region_East North Central float64
33 us_region_Mid-Atlantic float64
34 us_region_Pacific float64
35 us_region_South Atlantic float64
dtypes: float64(35), int64(1)
memory usage: 1.2 GB
```

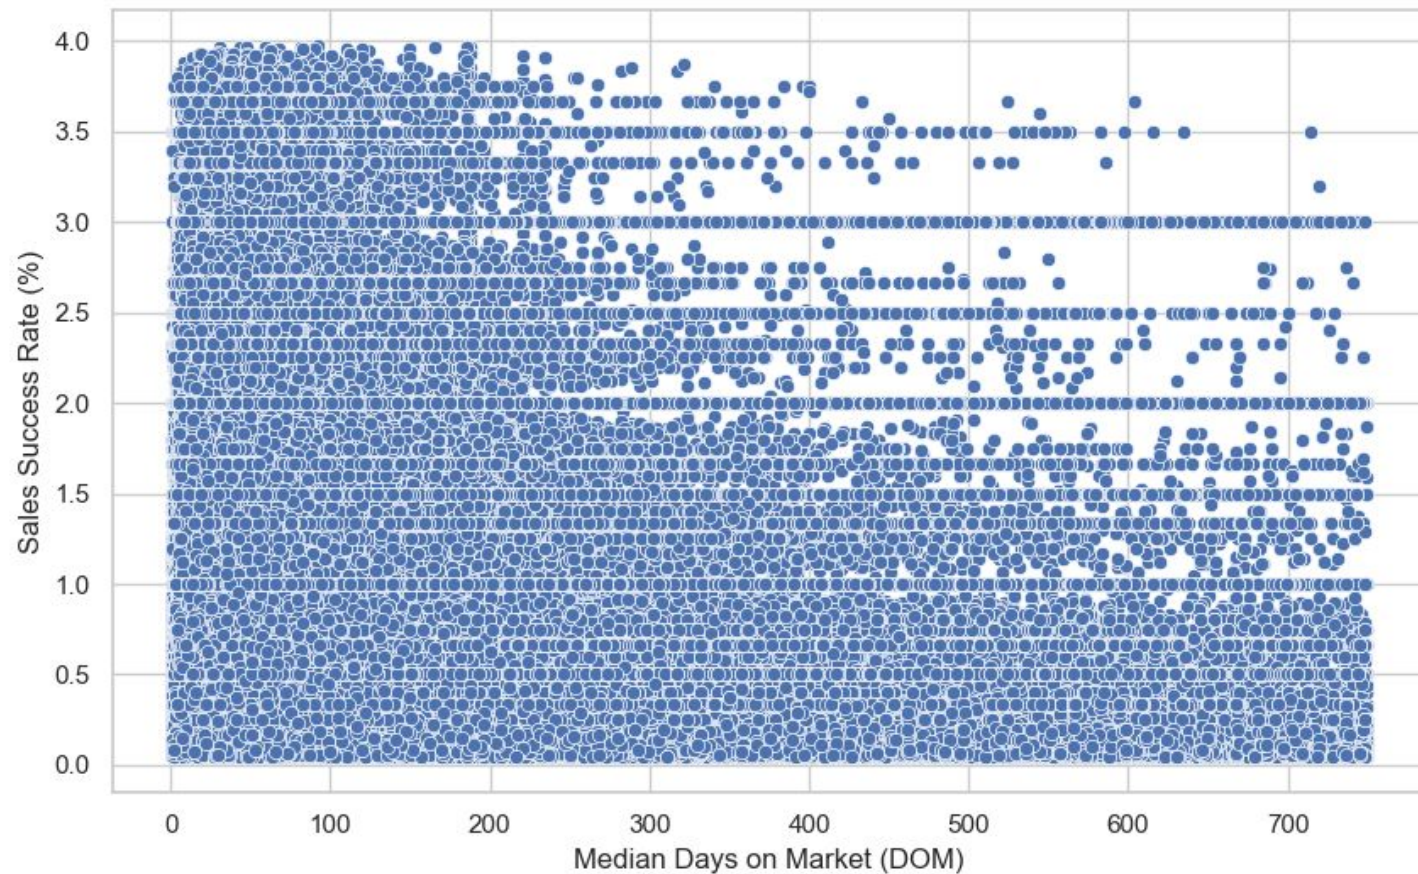
Preprocessing

- Nulls
- Combining Features
- Duplicates
- Adjusted Prices
- Binning by Regions
- Outliers, Irrelevant Features
- Highly correlated features

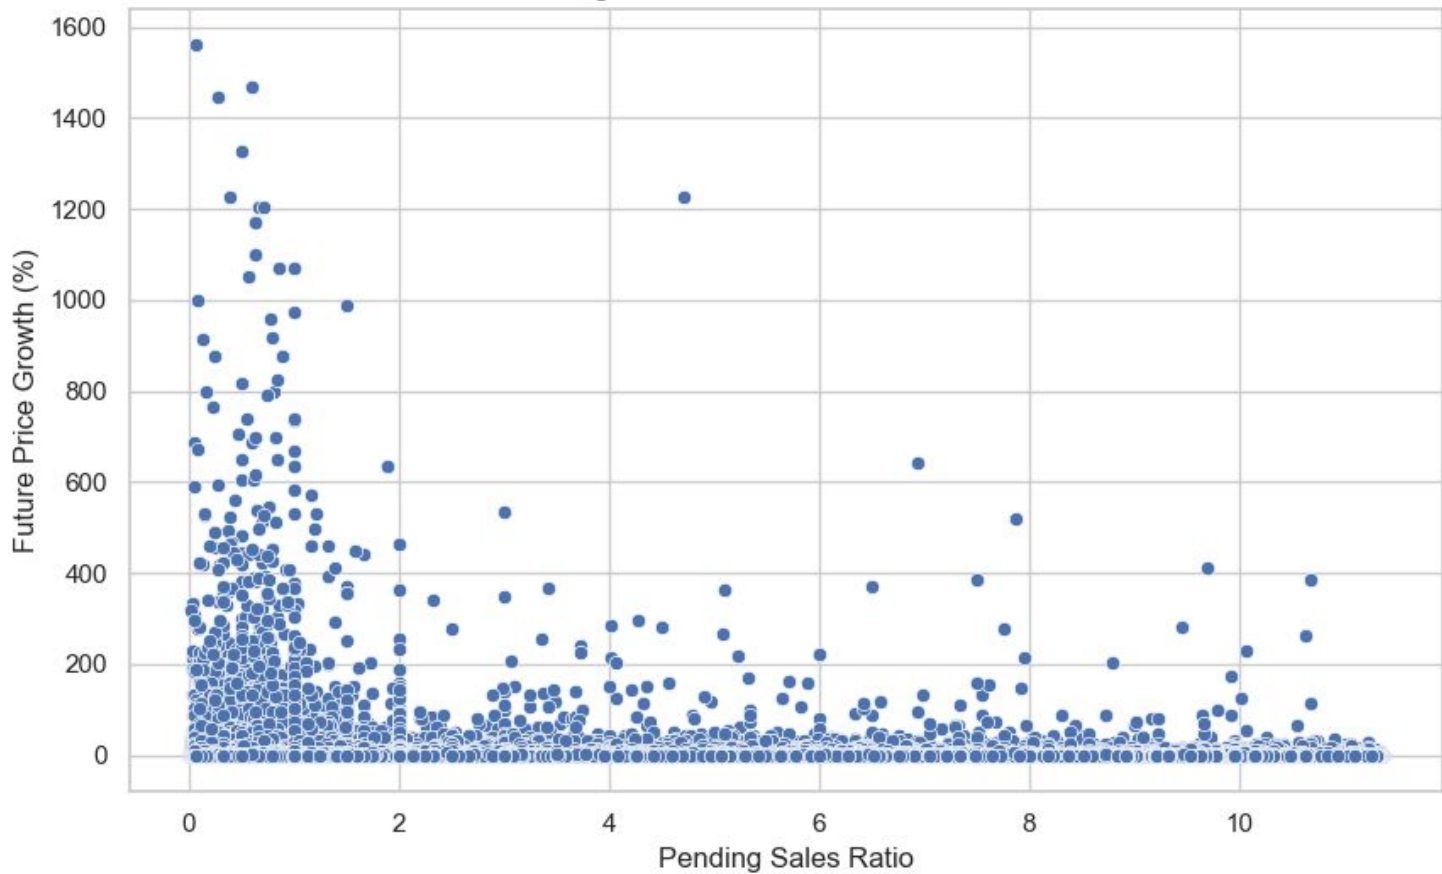
Correlation Matrix of Variables



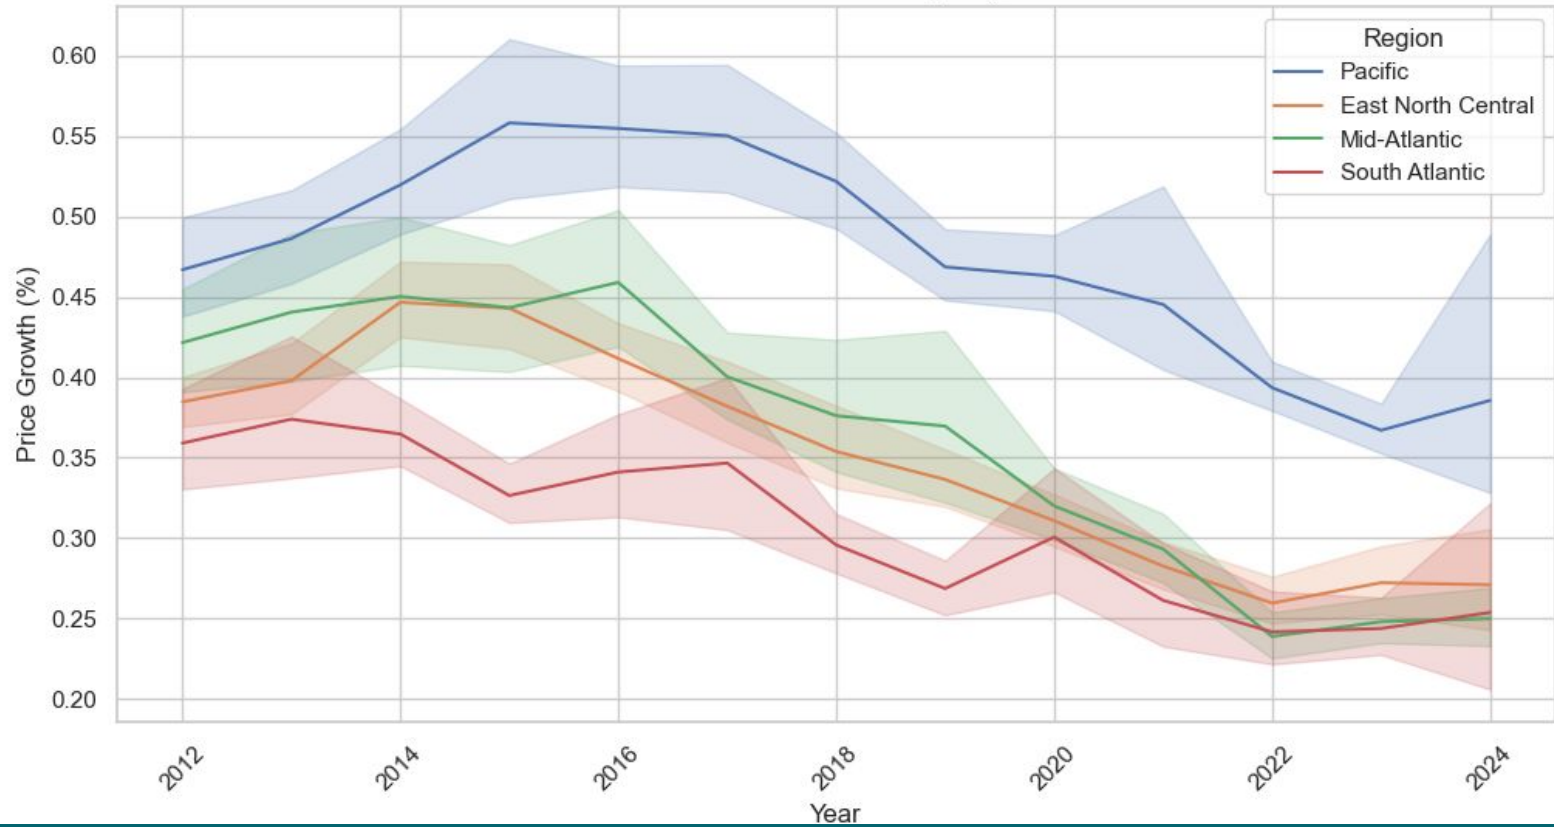
Median DOM vs. Sales Success Rate



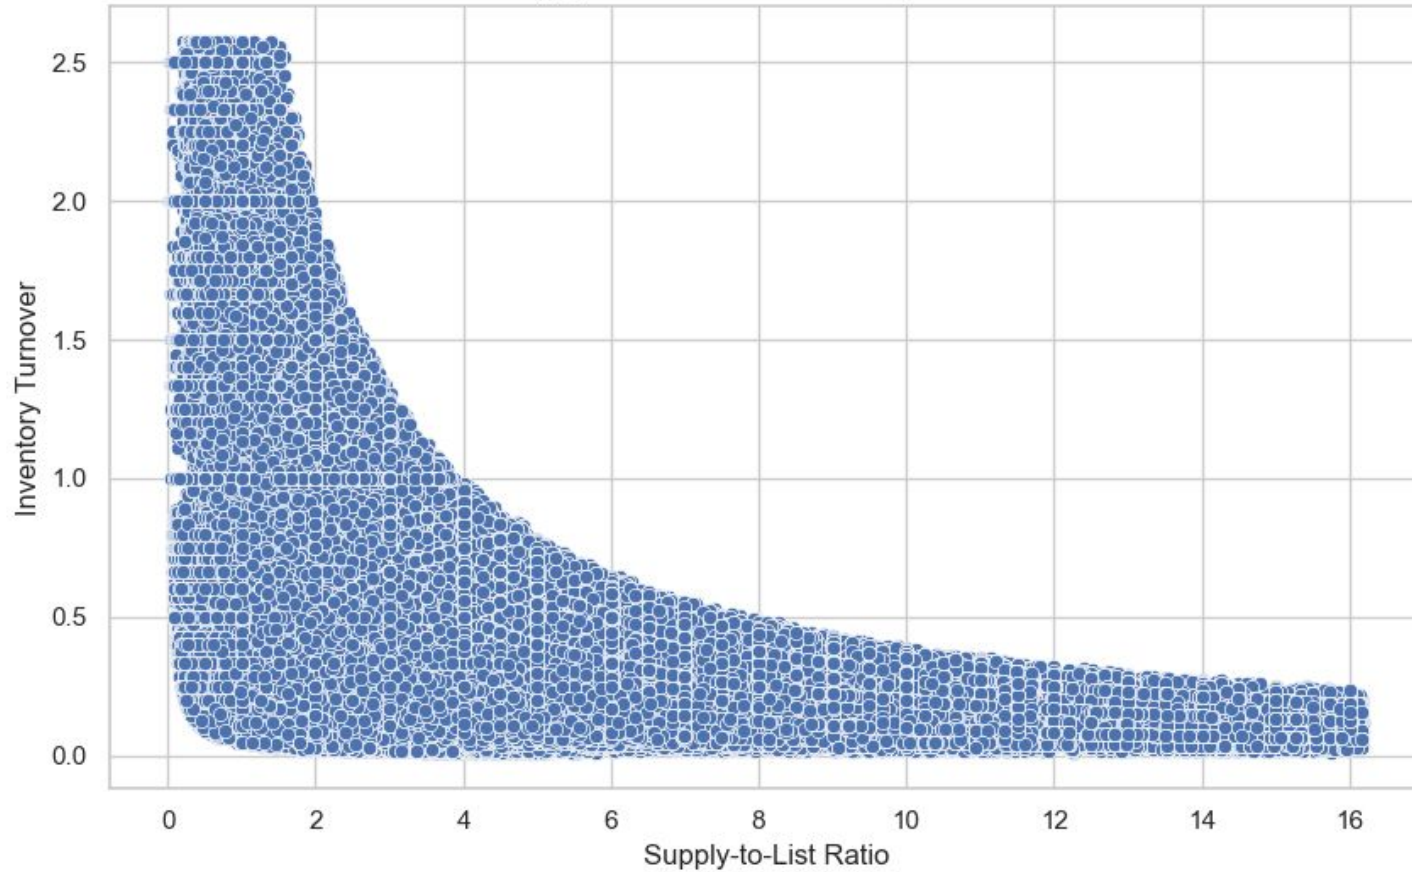
Pending Sales Ratio vs. Future Price Growth

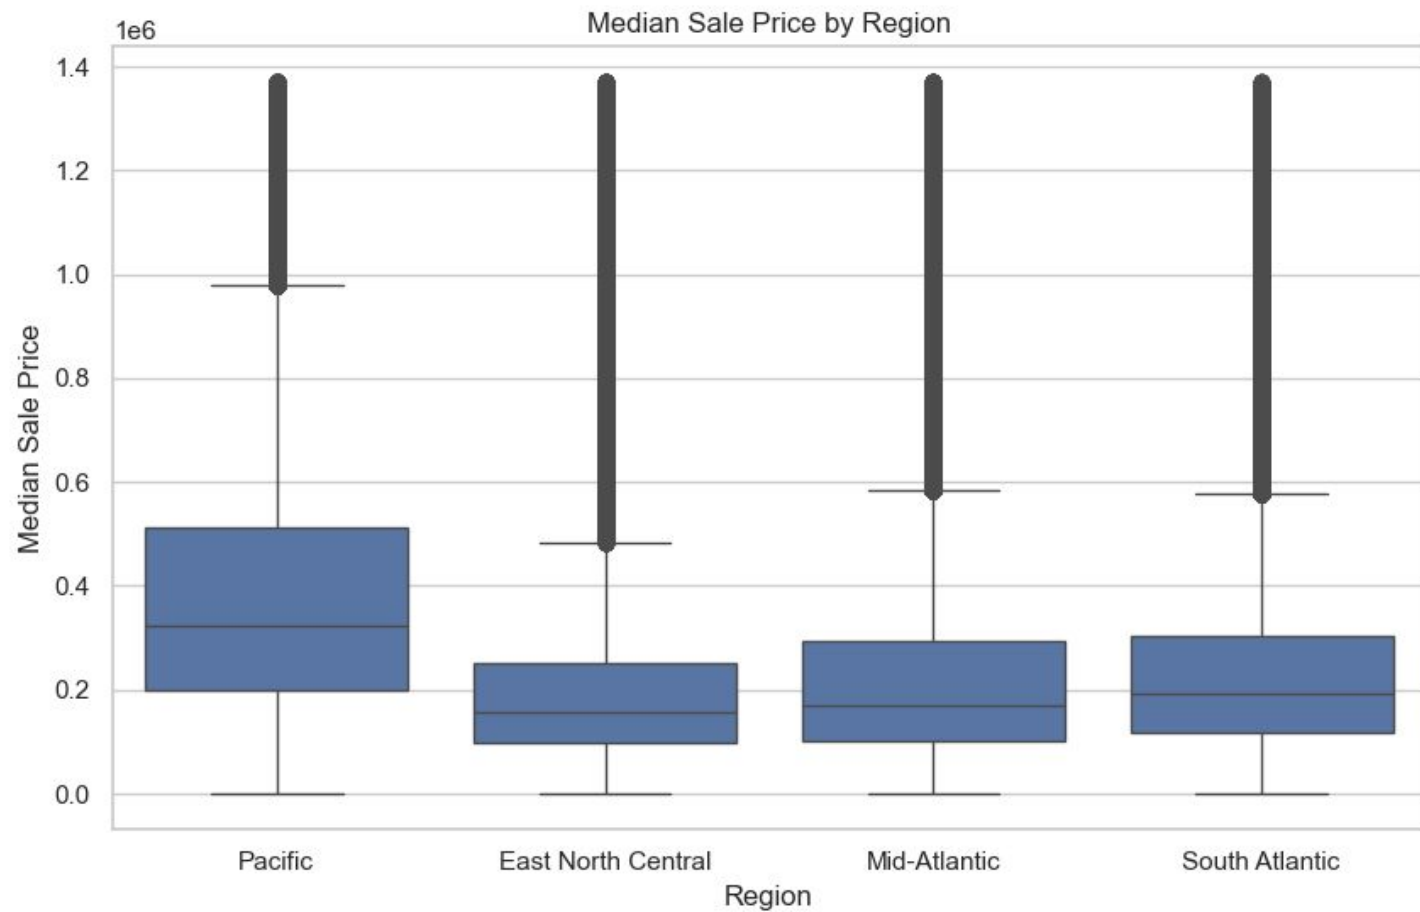


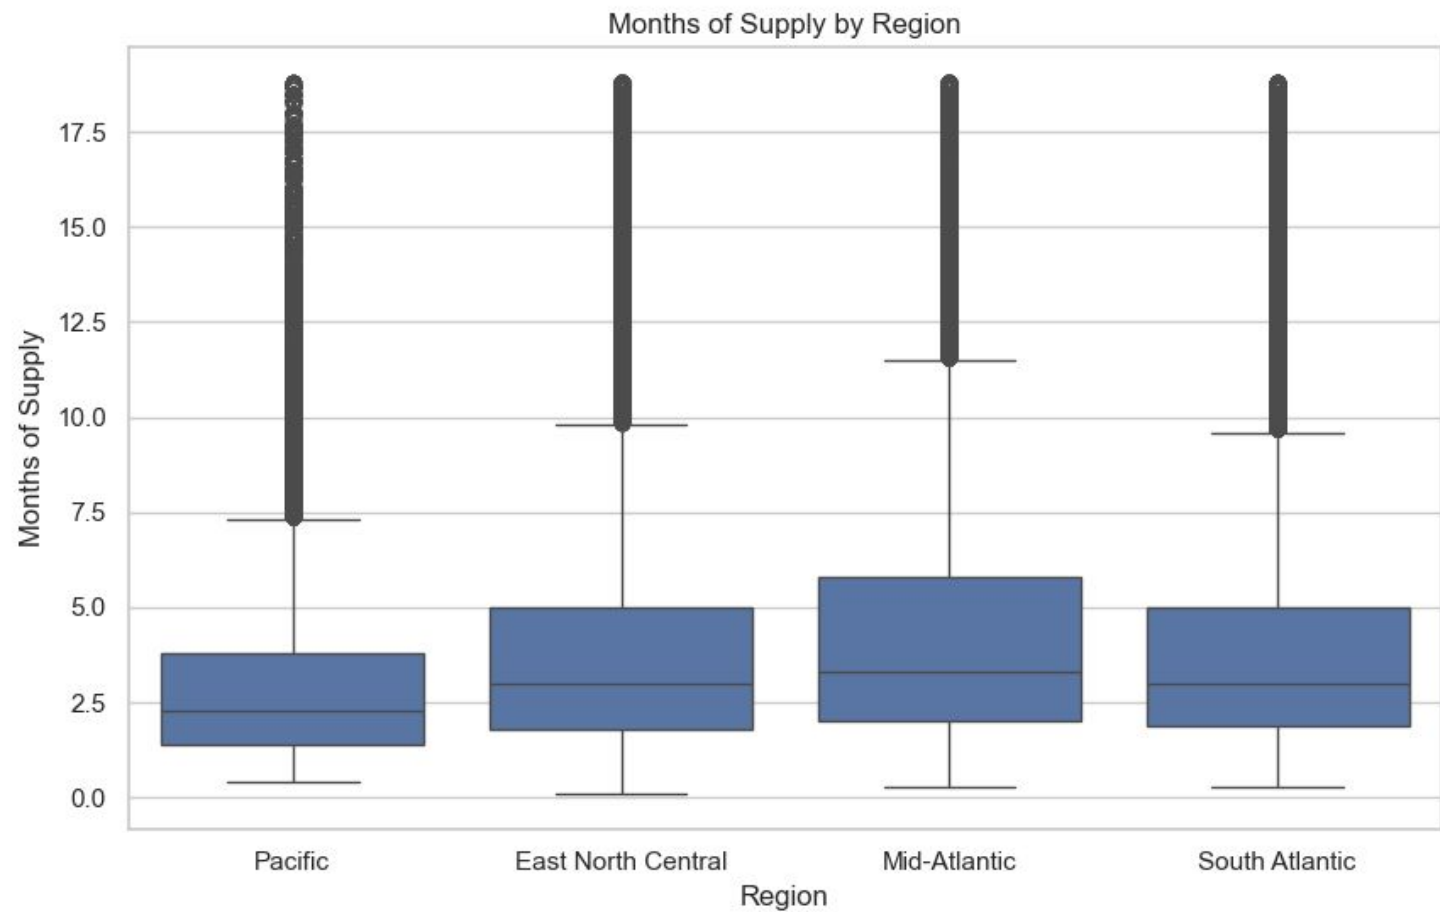
Price Growth Over Time by Region



Supply-to-List Ratio vs. Inventory Turnover







```
# Feature selection - selecting columns that are likely to affect homes sold
features = ['median_sale_price', 'median_list_price', 'inventory', 'months_of_supply',
            'price_growth', 'buyer_utility', 'pending_sales_ratio', 'sales_success_rate',
            'inventory_turnover', 'adjusted_months_supply', 'supply_to_list_ratio',
            'property_type_All Residential', 'state_avg_sale_price', 'price_momentum',
            'supply_pressure', 'demand_pressure', 'price_elasticity']
```

```
X = df[features] # Predictor variables
y = df['homes_sold'] # Target variable
```

```
# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Initialize and train the linear regression model
lr_model = LinearRegression()
lr_model.fit(X_train_scaled, y_train)
```

```
# Predict homes sold using the test set
y_pred_lr = lr_model.predict(X_test_scaled)
```

```
# Evaluate the model
mse_lr = mean_squared_error(y_test, y_pred_lr)
r2_lr = r2_score(y_test, y_pred_lr)
```

```
print(f'Linear Regression - MSE: {mse_lr}, R-squared: {r2_lr}')
```

```
Linear Regression - MSE: 100.73826222540995, R-squared: 0.755139756428168
```

```
X = df.drop(columns=['median_sale_price', 'Unnamed: 0']) # Predictor variables
y = df['median_sale_price'] # Target variable

# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Initialize the Linear Regression model
lr_model = LinearRegression()
```

```
# Fit the model to the training data
lr_model.fit(X_train, y_train)
```

```
# Predict buyer utility on the test data
y_pred = lr_model.predict(X_test)
```

```
# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse}')
```

```
# Calculate R-squared ( $R^2$ )
r2 = r2_score(y_test, y_pred)
print(f'R-squared: {r2}')
```

```
Mean Squared Error: 5983453454.910419
R-squared: 0.847707192781659
```

- Increase accuracy of models
- Random Forest?
- KNN clustering?



